

## Chapter 9.7a Using JAGS

Jim Albert and Monika Hu

Chapter 9 Simulation by Markov Chain Monte Carlo

# Introduction

- ▶ There has been an effort to develop general-purpose Bayesian computing software.
- ▶ One of the earliest Bayesian simulation-based computing software was BUGS (for Bayesian inference Using Gibbs Sampling) and we illustrate a similar package JAGS (for Just Another Gibbs Sampler).
- ▶ In using JAGS, one defines a Bayesian model by writing a short script. One then inputs this script together with data and prior parameter values in a single R function from the `runjags` package. This function simulates from the MCMC algorithm for a specified number of samples.

## Normal sampling model

- ▶ Consider the problem of estimating the mean Buffalo snowfall assuming Normal sampling with both the mean and standard deviation unknown and independent priors.
- ▶ Express the parameters of the Normal distribution as  $\mu$  and the precision  $\phi$ .
- ▶ Write this Bayesian model as
- ▶ Sampling, for  $i = 1, \dots, n$ :

$$Y_i \stackrel{i.i.d.}{\sim} \text{Normal}(\mu, \sqrt{1/\phi}).$$

- ▶ Independent priors for  $\mu$  and  $\phi$ :

$$\mu \sim \text{Normal}(\mu_0, \sqrt{1/\phi_0}), \phi \sim \text{Gamma}(a, b).$$

# Describe the model by a script

- Write the following script defining this model.

```
modelString = "  
model{  
  ## sampling  
  for (i in 1:N) {  
    y[i] ~ dnorm(mu, phi)  
  }  
  ## priors  
  mu ~ dnorm(mu0, phi0)  
  phi ~ dgamma(a, b)  
  sigma <- sqrt(pow(phi, -1))  
}
```

# Comments

- ▶ Note that this script resembles the statement of the model.
- ▶ In the sampling part, the loop structure `for (i in 1:N)` is used to assign the distribution of each value in the data vector `y` the same Normal distribution, represented by `dnorm`.
- ▶ In the prior part , use `dnorm` and `dgamma` to specify the Normal prior and Gamma prior for `mu` and `phi`.

## Define the data and prior parameters

- In the script below, a list `the_data` is used to collect the observations `y`, the number of observations `N`, and values of the Normal prior parameters `mu0`, `phi0`, and of the Gamma prior parameters `a` and `b`.

```
buffalo <- read.csv("../data/buffalo_snowfall.csv")
data <- buffalo[59:78, c("SEASON", "JAN")]
y <- data$JAN
N <- length(y)
the_data <- list("y" = y, "N" = N,
                 "mu0"=10, "phi0"=1/3^2,
                 "a"=1, "b"=1)
```

## Define initial values

- ▶ One supplies initial values in the MCMC simulation for all of the parameters.
- ▶ Alternatively, one can specify the initial values by means of a function – this will be implemented when multiple chains are discussed.
- ▶ If no initial values are specified, then JAGS will select initial values – these are usually a “typical” value such as a mean or median from the prior distribution.

## Generate samples from the posterior

- ▶ The `run.jags()` from the `runjags` package does the sampling.
- ▶ The input `n.chains = 1` indicates that one stream of simulated values will be generated.
- ▶ `adapt = 1000` says that 1000 simulated iterations are used in “adapt period” to prepare for MCMC
- ▶ `burnin = 1000` indicates 5000 simulated iterations are used in a “burn-in” period
- ▶ `sample = 5000` arguments indicates that 5000 iterations of the algorithm will be collected.
- ▶ The `monitor` arguments says that we are collecting simulated values of the mean `mu` and the standard deviation `sigma`.



## run.jags() function

```
posterior <- run.jags(modelString,  
                      n.chains = 1,  
                      data = the_data,  
                      monitor = c("mu", "sigma"),  
                      adapt = 1000,  
                      burnin = 5000,  
                      sample = 5000,  
                      inits = initsfunction)
```

# MCMC diagnostics and summarization

- ▶ Before summarizing the simulated sample, some graphical diagnostics methods should be implemented to judge if the sample appears to move well across the space of likely values of the parameters.
- ▶ The `plot()` function in the `runjags` package constructs a collection of four graphs for a parameter of interest.
- ▶ By running `plot()` for `mu` and `sigma`, we obtain the graphs displayed in the next slide.

```
plot(posterior, vars = "mu")  
plot(posterior, vars = "sigma")
```

# Diagnostic Plots for $\mu$

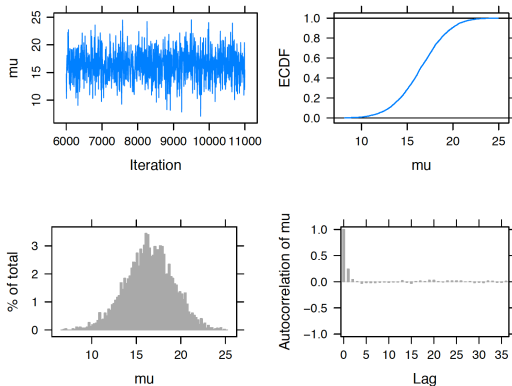


Figure 1: Diagnostic plots of simulated draws of mean using the JAGS software with the runjags package.

# Diagnostic Plots for $\sigma$

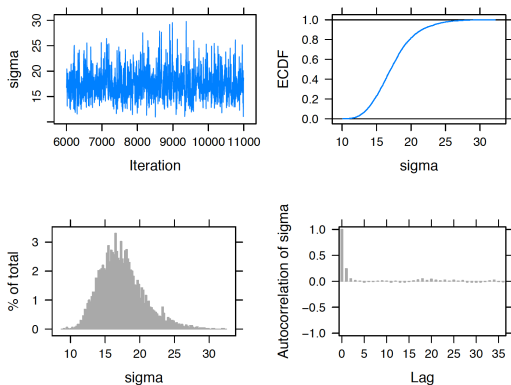


Figure 2: Diagnostic plots of simulated draws of standard deviation using the JAGS software with the runjags package.

# Diagnostic plots

- ▶ Trace and autocorrelation plots are helpful for seeing how the sampler moves across the posterior distribution.
- ▶ Here the trace plots show little autocorrelation and both simulated samples of  $\mu$  and  $\sigma$  appear to mix well.
- ▶ In the autocorrelation plots, the value of the autocorrelation drops sharply to zero as a function of the lag which confirms modest autocorrelation.

= The remaining graphs are a histogram of the simulated draws and an estimate at the cumulative distribution function.

# Posterior Summaries

- ▶ We are encouraged by these diagnostic graphs so we obtain summaries of the simulated samples of  $\mu$  and  $\sigma$
- ▶ Use `print()` function on our MCMC object.
- ▶ The posterior mean of  $\mu$  is 16.5.
- ▶ The standard error of this simulation estimate is the “MCerr” value of 0.0486 – this standard error takes in account the correlated nature of these simulated draws.

```
print(posterior, digits = 3)
```

	Lower95	Median	Upper95	Mean	SD	Mode	MCerr
mu	10.8	16.5	21.4	16.5	2.68	--	0.0486
sigma	11.8	17.1	24	17.4	3.18	--	0.0576

## Multiple chains

- ▶ Can try different starting values and running several MCMC chains.
- ▶ This is facilitated by arguments in the `run.jags()` function. Suppose one considers the very different pairs of starting values,  $(\mu, \phi) = (2, 1/4)$  and  $(\mu, \phi) = (30, 1/900)$ .
- ▶ Define a value `InitialValues`, a list containing two lists, each list containing a starting value.

```
InitialValues <- list(  
  list(mu = 2, phi = 1 / 4),  
  list(mu = 30, phi = 1 / 900)  
)
```

# Multiple Chains

- ▶ The `run.jags()` function is run with two modifications
- ▶ One chooses `n.chains = 2` and the initial values are input through the `inits = InitialValues` option.

```
posterior <- run.jags(modelString,  
                      n.chains = 2,  
                      data = the_data,  
                      monitor = c("mu", "sigma"),  
                      adapt = 1000,  
                      burnin = 5000,  
                      sample = 5000,  
                      inits = InitialValues)
```



# Output

- ▶ The output variable posterior contains simulated draws from both chains.
- ▶ One compares posterior quantiles from the two chains. – they are close in value indicating that the MCMC run is insensitive to the choice of starting value.

```
summary(posterior$mcmc[[1]], digits = 3)
```

	2.5%	25%	50%	75%	97.5%
mu	10.99	14.64	16.49	18.35	21.62
sigma	12.26	15.15	17.03	19.31	25.07

```
summary(posterior$mcmc[[2]], digits = 3)
```

	2.5%	25%	50%	75%	97.5%
mu	10.97	14.59	16.55	18.33	21.54
sigma	12.21	15.08	16.96	19.18	24.99